

# 02-DataFrames

October 19, 2022

---

Copyright Pierian Data

For more information, visit us at [www.pieriandata.com](http://www.pieriandata.com)

## 1 DataFrames

DataFrames are the workhorse of pandas and are directly inspired by the R programming language. We can think of a DataFrame as a bunch of Series objects put together to share the same index. Let's use pandas to explore this topic!

```
[1]: import pandas as pd
import numpy as np

[2]: from numpy.random import randn
np.random.seed(101)

[3]: df = pd.DataFrame(randn(5,4),index='A B C D E'.split(),columns='W X Y Z'.
↳split())

[4]: df
```

```
[4]:
```

|   | W         | X         | Y         | Z         |
|---|-----------|-----------|-----------|-----------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

### 1.1 Selection and Indexing

Let's learn the various methods to grab data from a DataFrame

```
[5]: df['W']
```

```
[5]: A    2.706850
      B    0.651118
      C   -2.018168
      D    0.188695
      E    0.190794
      Name: W, dtype: float64
```

```
[6]: # Pass a list of column names
      df[['W', 'Z']]
```

```
[6]:           W           Z
A    2.706850    0.503826
B    0.651118    0.605965
C   -2.018168   -0.589001
D    0.188695    0.955057
E    0.190794    0.683509
```

```
[7]: # SQL Syntax (NOT RECOMMENDED!)
      df.W
```

```
[7]: A    2.706850
      B    0.651118
      C   -2.018168
      D    0.188695
      E    0.190794
      Name: W, dtype: float64
```

DataFrame Columns are just Series

```
[8]: type(df['W'])
```

```
[8]: pandas.core.series.Series
```

### 1.1.1 Creating a new column:

```
[9]: df['new'] = df['W'] + df['Y']
```

```
[10]: df
```

```
[10]:           W           X           Y           Z          new
A    2.706850    0.628133    0.907969    0.503826    3.614819
B    0.651118   -0.319318   -0.848077    0.605965   -0.196959
C   -2.018168    0.740122    0.528813   -0.589001   -1.489355
D    0.188695   -0.758872   -0.933237    0.955057   -0.744542
E    0.190794    1.978757    2.605967    0.683509    2.796762
```

### 1.1.2 Removing Columns

```
[11]: df.drop('new',axis=1)
```

```
[11]:
```

|   | W         | X         | Y         | Z         |
|---|-----------|-----------|-----------|-----------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

```
[12]: # Not inplace unless specified!  
df
```

```
[12]:
```

|   | W         | X         | Y         | Z         | new       |
|---|-----------|-----------|-----------|-----------|-----------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  | 3.614819  |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  | -0.196959 |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 | -1.489355 |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  | -0.744542 |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  | 2.796762  |

```
[13]: df.drop('new',axis=1,inplace=True)
```

```
[14]: df
```

```
[14]:
```

|   | W         | X         | Y         | Z         |
|---|-----------|-----------|-----------|-----------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

Can also drop rows this way:

```
[15]: df.drop('E',axis=0)
```

```
[15]:
```

|   | W         | X         | Y         | Z         |
|---|-----------|-----------|-----------|-----------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  |

### 1.1.3 Selecting Rows

```
[16]: df.loc['A']
```

```
[16]: W    2.706850
      X    0.628133
      Y    0.907969
      Z    0.503826
      Name: A, dtype: float64
```

Or select based off of position instead of label

```
[17]: df.iloc[2]
```

```
[17]: W    -2.018168
      X    0.740122
      Y    0.528813
      Z   -0.589001
      Name: C, dtype: float64
```

#### 1.1.4 Selecting subset of rows and columns

```
[18]: df.loc['B', 'Y']
```

```
[18]: -0.8480769834036315
```

```
[19]: df.loc[['A', 'B'], ['W', 'Y']]
```

```
[19]:      W      Y
A  2.706850  0.907969
B  0.651118 -0.848077
```

#### 1.1.5 Conditional Selection

An important feature of pandas is conditional selection using bracket notation, very similar to numpy:

```
[20]: df
```

```
[20]:      W      X      Y      Z
A  2.706850  0.628133  0.907969  0.503826
B  0.651118 -0.319318 -0.848077  0.605965
C -2.018168  0.740122  0.528813 -0.589001
D  0.188695 -0.758872 -0.933237  0.955057
E  0.190794  1.978757  2.605967  0.683509
```

```
[21]: df>0
```

```
[21]:
```

|   | W     | X     | Y     | Z     |
|---|-------|-------|-------|-------|
| A | True  | True  | True  | True  |
| B | True  | False | False | True  |
| C | False | True  | True  | False |
| D | True  | False | False | True  |
| E | True  | True  | True  | True  |

```
[22]: df[df>0]
```

```
[22]:
```

|   | W        | X        | Y        | Z        |
|---|----------|----------|----------|----------|
| A | 2.706850 | 0.628133 | 0.907969 | 0.503826 |
| B | 0.651118 | NaN      | NaN      | 0.605965 |
| C | NaN      | 0.740122 | 0.528813 | NaN      |
| D | 0.188695 | NaN      | NaN      | 0.955057 |
| E | 0.190794 | 1.978757 | 2.605967 | 0.683509 |

```
[23]: df[df['W']>0]
```

```
[23]:
```

|   | W        | X         | Y         | Z        |
|---|----------|-----------|-----------|----------|
| A | 2.706850 | 0.628133  | 0.907969  | 0.503826 |
| B | 0.651118 | -0.319318 | -0.848077 | 0.605965 |
| D | 0.188695 | -0.758872 | -0.933237 | 0.955057 |
| E | 0.190794 | 1.978757  | 2.605967  | 0.683509 |

```
[24]: df[df['W']>0]['Y']
```

```
[24]: A    0.907969
      B   -0.848077
      D   -0.933237
      E    2.605967
      Name: Y, dtype: float64
```

```
[25]: df[df['W']>0][['Y','X']]
```

```
[25]:
```

|   | Y         | X         |
|---|-----------|-----------|
| A | 0.907969  | 0.628133  |
| B | -0.848077 | -0.319318 |
| D | -0.933237 | -0.758872 |
| E | 2.605967  | 1.978757  |

For two conditions you can use | and & with parenthesis:

```
[26]: df[(df['W']>0) & (df['Y'] > 1)]
```

```
[26]:
```

|   | W        | X        | Y        | Z        |
|---|----------|----------|----------|----------|
| E | 0.190794 | 1.978757 | 2.605967 | 0.683509 |

## 1.2 More Index Details

Let's discuss some more features of indexing, including resetting the index or setting it something else. We'll also talk about index hierarchy!

```
[27]: df
```

```
[27]:
```

|   | W         | X         | Y         | Z         |
|---|-----------|-----------|-----------|-----------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

```
[28]: # Reset to default 0,1...n index
df.reset_index()
```

```
[28]:
```

|   | index | W         | X         | Y         | Z         |
|---|-------|-----------|-----------|-----------|-----------|
| 0 | A     | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| 1 | B     | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| 2 | C     | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| 3 | D     | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| 4 | E     | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

```
[29]: newind = 'CA NY WY OR CO'.split()
```

```
[30]: df['States'] = newind
```

```
[31]: df
```

```
[31]:
```

|   | W         | X         | Y         | Z         | States |
|---|-----------|-----------|-----------|-----------|--------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  | CA     |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  | NY     |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 | WY     |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  | OR     |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  | CO     |

```
[32]: df.set_index('States')
```

```
[32]:
```

|        | W         | X         | Y         | Z         |
|--------|-----------|-----------|-----------|-----------|
| States |           |           |           |           |
| CA     | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| NY     | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| WY     | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| OR     | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| CO     | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

```
[33]: df
```

```
[33]:
```

|   | W         | X         | Y         | Z         | States |
|---|-----------|-----------|-----------|-----------|--------|
| A | 2.706850  | 0.628133  | 0.907969  | 0.503826  | CA     |
| B | 0.651118  | -0.319318 | -0.848077 | 0.605965  | NY     |
| C | -2.018168 | 0.740122  | 0.528813  | -0.589001 | WY     |
| D | 0.188695  | -0.758872 | -0.933237 | 0.955057  | OR     |
| E | 0.190794  | 1.978757  | 2.605967  | 0.683509  | CO     |

```
[34]: df.set_index('States',inplace=True)
```

```
[35]: df
```

```
[35]:
```

|        | W         | X         | Y         | Z         |
|--------|-----------|-----------|-----------|-----------|
| States |           |           |           |           |
| CA     | 2.706850  | 0.628133  | 0.907969  | 0.503826  |
| NY     | 0.651118  | -0.319318 | -0.848077 | 0.605965  |
| WY     | -2.018168 | 0.740122  | 0.528813  | -0.589001 |
| OR     | 0.188695  | -0.758872 | -0.933237 | 0.955057  |
| CO     | 0.190794  | 1.978757  | 2.605967  | 0.683509  |

### 1.3 DataFrame Summaries

There are a couple of ways to obtain summary data on DataFrames. `df.describe()` provides summary statistics on all numerical columns. `df.info` and `df.dtypes` displays the data type of all columns.

```
[36]: df.describe()
```

```
[36]:
```

|       | W         | X         | Y         | Z         |
|-------|-----------|-----------|-----------|-----------|
| count | 5.000000  | 5.000000  | 5.000000  | 5.000000  |
| mean  | 0.343858  | 0.453764  | 0.452287  | 0.431871  |
| std   | 1.681131  | 1.061385  | 1.454516  | 0.594708  |
| min   | -2.018168 | -0.758872 | -0.933237 | -0.589001 |
| 25%   | 0.188695  | -0.319318 | -0.848077 | 0.503826  |
| 50%   | 0.190794  | 0.628133  | 0.528813  | 0.605965  |
| 75%   | 0.651118  | 0.740122  | 0.907969  | 0.683509  |
| max   | 2.706850  | 1.978757  | 2.605967  | 0.955057  |

```
[38]: df.dtypes
```

```
[38]: W    float64  
X    float64  
Y    float64  
Z    float64  
dtype: object
```

```
[40]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 5 entries, CA to CO  
Data columns (total 4 columns):  
W      5 non-null float64  
X      5 non-null float64  
Y      5 non-null float64  
Z      5 non-null float64  
dtypes: float64(4)  
memory usage: 200.0+ bytes
```

## 2 Great Job!