

01-ACF-and-PACF

October 19, 2022

Copyright Pierian Data

For more information, visit us at www.pieriandata.com

1 ACF and PACF

2 Autocorrelation Function / Partial Autocorrelation Function

Before we can investigate autoregression as a modeling tool, we need to look at covariance and correlation as they relate to lagged (shifted) samples of a time series.

2.0.1 Goals

- Be able to create ACF and PACF charts
- Create these charts for multiple times series, one with seasonality and another without
- Be able to calculate Orders PQD terms for ARIMA off these charts (highlight where they cross the x axis)

Related Functions:

`stattools.acovf(x[, unbiased, demean, fft, ...])` Autocovariance for 1D
`stattools.acf(x[, unbiased, nlags, qstat, ...])` Autocorrelation function for 1d arrays
`stattools.pacf(x[, nlags, method, alpha])` Partial autocorrelation estimated
`stattools.pacf_yw(x[, nlags, method])` Partial autocorrelation estimated with non-recursive yule_walker
`stattools.pacf_ols(x[, nlags])` Calculate partial autocorrelations

Related Plot Methods:

`tsaplots.plot_acf(x)` Plot the autocorrelation function
`tsaplots.plot_pacf(x)` Plot the partial autocorrelation function

For Further Reading:

Wikipedia: Autocovariance Forecasting: Principles and Practice Autocorrelation NIST Statistics Handbook Partial Autocorrelation Plot

2.1 Perform standard imports and load datasets

```
[1]: import pandas as pd
import numpy as np
%matplotlib inline
import statsmodels.api as sm

# Load a non-stationary dataset
df1 = pd.read_csv('../Data/airline_passengers.
    ↪csv', index_col='Month', parse_dates=True)
df1.index.freq = 'MS'

# Load a stationary dataset
df2 = pd.read_csv('../Data/DailyTotalFemaleBirths.
    ↪csv', index_col='Date', parse_dates=True)
df2.index.freq = 'D'

[2]: # Import the models we'll be using in this section
from statsmodels.tsa.stattools import acovf, acf, pacf, pacf_yw, pacf_ols
```

2.1.1 Ignore harmless warnings

A quick note before we get started. Many of the models used in this and upcoming sections are likely to raise harmless warnings. For instance, the unbiased partial autocorrelation `pacf_yw()` performed below may raise a `RuntimeWarning`: invalid value encountered in `sqrt`. We don't really need to be concerned with this, and we can avoid it with the following code:

```
[3]: import warnings
warnings.filterwarnings("ignore")
```

2.2 Autocovariance for 1D

In a deterministic process, like $y = \sin(x)$, we always know the value of y for a given value of x . However, in a stochastic process there is always some randomness that prevents us from knowing the value of y . Instead, we analyze the past (or lagged) behavior of the system to derive a probabilistic estimate for \hat{y} .

One useful descriptor is covariance. When talking about dependent and independent x and y variables, covariance describes how the variance in x relates to the variance in y . Here the size of the covariance isn't really important, as x and y may have very different scales. However, if the covariance is positive it means that x and y are changing in the same direction, and may be related.

With a time series, x is a fixed interval. Here we want to look at the variance of y_t against lagged or shifted values of y_{t+k}

For a stationary time series, the autocovariance function for γ (gamma) is given as:

$$\gamma_{XX}(t_1, t_2) = \text{Cov}[X_{t_1}, X_{t_2}] = E[(X_{t_1} - \mu_{t_1})(X_{t_2} - \mu_{t_2})]$$

We can calculate a specific γ_k with:

$$\gamma_k = \frac{1}{n} \sum_{t=1}^{n-k} (y_t - \bar{y})(y_{t+k} - \bar{y})$$

A NOTE ON FORMULA CONVENTIONS: Different texts employ different symbol conventions. For example, in the above autocovariance formula we use k to represent the amount of lag or shift. Some texts use h instead.

2.2.1 Autocovariance Example:

Say we have a time series with five observations: $\{13, 5, 11, 12, 9\}$. We can quickly see that $n = 5$, the mean $\bar{y} = 10$, and we'll see that the variance $\sigma^2 = 8$. The following calculations give us our covariance values: $\gamma_0 = \frac{(13-10)(13-10)+(5-10)(5-10)+(11-10)(11-10)+(12-10)(12-10)+(9-10)(9-10)}{5} = \frac{40}{5} = 8.0$

$$\gamma_1 = \frac{(13-10)(5-10)+(5-10)(11-10)+(11-10)(12-10)+(12-10)(9-10)}{5} = \frac{-20}{5} = -4.0$$

$$\gamma_2 = \frac{(13-10)(11-10)+(5-10)(12-10)+(11-10)(9-10)}{5} = \frac{-8}{5} = -1.6$$

$$\gamma_3 = \frac{(13-10)(12-10)+(5-10)(9-10)}{5} = \frac{11}{5} = 2.2$$

$$\gamma_4 = \frac{(13-10)(9-10)}{5} = \frac{-3}{5} = -0.6 \text{ Note that } \gamma_0 \text{ is just the population variance } \sigma^2$$

Let's see if statsmodels gives us the same results! For this we'll create a fake dataset:

```
[4]: df = pd.DataFrame({'a':[13, 5, 11, 12, 9]})
      arr = acovf(df['a'])
      arr
```

```
[4]: array([ 8. , -4. , -1.6,  2.2, -0.6])
```

2.2.2 Unbiased Autocovariance

Note that the number of terms in the calculations above are decreasing. Statsmodels can return an "unbiased" autocovariance where instead of dividing by n we divide by $n - k$.

$$\gamma_0 = \frac{(13-10)(13-10)+(5-10)(5-10)+(11-10)(11-10)+(12-10)(12-10)+(9-10)(9-10)}{5-0} = \frac{40}{5} = 8.0$$

$$\gamma_1 = \frac{(13-10)(5-10)+(5-10)(11-10)+(11-10)(12-10)+(12-10)(9-10)}{5-1} = \frac{-20}{4} = -5.0$$

$$\gamma_2 = \frac{(13-10)(11-10)+(5-10)(12-10)+(11-10)(9-10)}{5-2} = \frac{-8}{3} = -2.67$$

$$\gamma_3 = \frac{(13-10)(12-10)+(5-10)(9-10)}{5-3} = \frac{11}{2} = 5.5$$

$$\gamma_4 = \frac{(13-10)(9-10)}{5-4} = \frac{-3}{1} = -3.0$$

```
[5]: arr2 = acovf(df['a'],unbiased=True)
      arr2
```

```
[5]: array([ 8.          , -5.          , -2.66666667,  5.5          , -3.          ])
```

2.3 Autocorrelation for 1D

The correlation ρ (rho) between two variables y_1, y_2 is given as:

$$\mathbf{2.3.1} \quad \rho = \frac{E[(y_1 - \mu_1)(y_2 - \mu_2)]}{\sigma_1 \sigma_2} = \frac{\text{Cov}(y_1, y_2)}{\sigma_1 \sigma_2},$$

where E is the expectation operator, μ_1, σ_1 and μ_2, σ_2 are the means and standard deviations of y_1 and y_2 .

When working with a single variable (i.e. autocorrelation) we would consider y_1 to be the original series and y_2 a lagged version of it. Note that with autocorrelation we work with \bar{y} , that is, the full population mean, and not the means of the reduced set of lagged factors (see note below).

Thus, the formula for ρ_k for a time series at lag k is:

$$\rho_k = \frac{\sum_{t=1}^{n-k} (y_t - \bar{y})(y_{t+k} - \bar{y})}{\sum_{t=1}^n (y_t - \bar{y})^2}$$

This can be written in terms of the covariance constant γ_k as:

$$\rho_k = \frac{\gamma_k n}{\gamma_0 n} = \frac{\gamma_k}{\sigma^2}$$

For example, $\rho_4 = \frac{\gamma_4}{\sigma^2} = \frac{-0.6}{8} = -0.075$

Note that ACF values are bound by -1 and 1. That is, $-1 \leq \rho_k \leq 1$

```
[6]: arr3 = acf(df['a'])  
arr3
```

```
[6]: array([ 1.    , -0.5    , -0.2    ,  0.275, -0.075])
```

2.4 Partial Autocorrelation

Partial autocorrelations measure the linear dependence of one variable after removing the effect of other variable(s) that affect both variables. That is, the partial autocorrelation at lag k is the autocorrelation between y_t and y_{t+k} that is not accounted for by lags 1 through $k-1$.

A common method employs the non-recursive Yule-Walker Equations:

$$\begin{aligned}\phi_0 &= 1 \\ \phi_1 &= \rho_1 = -0.50 \\ \phi_2 &= \frac{\rho_2 - \rho_1^2}{1 - \rho_1^2} = \frac{(-0.20) - (-0.50)^2}{1 - (-0.50)^2} = \frac{-0.45}{0.75} = -0.60\end{aligned}$$

As k increases, we can solve for ϕ_k using matrix algebra and the Levinson–Durbin recursion algorithm which maps the sample autocorrelations ρ to a Toeplitz diagonal-constant matrix. The full solution is beyond the scope of this course, but the setup is as follows:

$$\begin{pmatrix} \rho_0 & \rho_1 & \cdots & \rho_{k-1} \\ \rho_1 & \rho_0 & \cdots & \rho_{k-2} \\ \vdots & \vdots & \ddots & \vdots \\ \rho_{k-1} & \rho_{k-2} & \cdots & \rho_0 \end{pmatrix} \begin{pmatrix} \phi_{k1} \\ \phi_{k2} \\ \vdots \\ \phi_{kk} \end{pmatrix} = \begin{pmatrix} \rho_1 \\ \rho_2 \\ \vdots \\ \rho_k \end{pmatrix}$$

```
[7]: arr4 = pacf_yw(df['a'],nlags=4,method='mle')
      arr4
```

```
[7]: array([ 1.          , -0.5          , -0.6          , -0.38541667, -0.40563273])
```

NOTE: We passed in method='mle' above in order to use biased ACF coefficients. "mle" stands for "maximum likelihood estimation". Alternatively we can pass method='unbiased' (the statsmodels default):

```
[8]: arr5 = pacf_yw(df['a'],nlags=4,method='unbiased')
      arr5
```

```
[8]: array([ 1.          , -0.625         , -1.18803419,  2.03764205,  0.8949589  ])
```

2.4.1 Partial Autocorrelation with OLS

This provides partial autocorrelations with ordinary least squares (OLS) estimates for each lag instead of Yule-Walker.

```
[9]: arr6 = pacf_ols(df['a'],nlags=4)
      arr6
```

```
[9]: array([ 1.          , -0.49677419, -0.43181818,  0.53082621,  0.25434783])
```


3 Plotting

The arrays returned by `.acf(df)` and `.pacf_yw(df)` show the magnitude of the autocorrelation for a given y at time t . Before we look at plotting arrays, let's look at the data itself for evidence of autocorrelation.

Pandas has a built-in plotting function that plots increasing y_t values on the horizontal axis against lagged versions of the values y_{t+1} on the vertical axis. If a dataset is non-stationary with an upward trend, then neighboring values should trend in the same way. Let's look at the Airline Passengers dataset first.

```
[10]: from pandas.plotting import lag_plot


      lag_plot(df1['Thousands of Passengers']);
```



Visually this shows evidence of a very strong autocorrelation; as y_t values increase, nearby (lagged) values also increase.

Now let's look at the stationary Daily Total Female Births dataset:

```
[11]: lag_plot(df2['Births']);
```



As expected, there is little evidence of autocorrelation here.

3.1 ACF Plots

Plotting the magnitude of the autocorrelations over the first few (20-40) lags can say a lot about a time series.


For example, consider the stationary Daily Total Female Births dataset:

```
[12]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
[13]: # Let's look first at the ACF array. By default acf() returns 40 lags
      acf(df2['Births'])
```

```
[13]: array([ 1.          ,  0.21724118,  0.15287758,  0.10821254,  0.09066059,
          0.09595481,  0.09104012,  0.19508071,  0.14115295,  0.06117859,
          0.04781522,  0.04770662, -0.01964707,  0.02287422,  0.08112657,
          0.11185686,  0.07333732,  0.01501845,  0.07270333,  0.06859   ,
          0.09280107,  0.26386846,  0.14012147,  0.06070286,  0.08716232,
          0.05038825,  0.0650489 ,  0.11466565,  0.1552232 ,  0.12850638,
          0.10358981,  0.09734643,  0.04912286,  0.04022798,  0.05838555,
          0.05359812,  0.10151053,  0.08268663,  0.0912185 ,  0.11192192,
          0.05652846])
```

```
[14]: # Now let's plot the autocorrelation at different lags
title = 'Autocorrelation: Daily Female Births'
lags = 40
plot_acf(df2,title=title,lags=lags);
```



This is a typical ACF plot for stationary data, with lags on the horizontal axis and correlations on the vertical axis. The first value y_0 is always 1. A sharp dropoff indicates that there is no AR component in the ARIMA model.

Next we'll look at non-stationary data with the Airline Passengers dataset:


```
[15]: acf(df1['Thousands of Passengers'])
```

```
[15]: array([1.          , 0.94804734, 0.87557484, 0.80668116, 0.75262542,
          0.71376997, 0.6817336 , 0.66290439, 0.65561048, 0.67094833,
          0.70271992, 0.74324019, 0.76039504, 0.71266087, 0.64634228,
          0.58592342, 0.53795519, 0.49974753, 0.46873401, 0.44987066,
          0.4416288 , 0.45722376, 0.48248203, 0.51712699, 0.53218983,
          0.49397569, 0.43772134, 0.3876029 , 0.34802503, 0.31498388,
          0.28849682, 0.27080187, 0.26429011, 0.27679934, 0.2985215 ,
          0.32558712, 0.3370236 , 0.30333486, 0.25397708, 0.21065534,
          0.17217092])
```

```
[16]: title = 'Autocorrelation: Airline Passengers'
lags = 40
```



```
plot_acf(df1,title=title,lags=lags);
```




This plot indicates non-stationary data, as there are a large number of lags before ACF values drop off.

3.2 PACF Plots

Partial autocorrelations work best with stationary data. Let's look first at Daily Total Female Births:


```
[17]: title='Partial Autocorrelation: Daily Female Births'  
      lags=40  
      plot_pacf(df2,title=title,lags=lags);
```



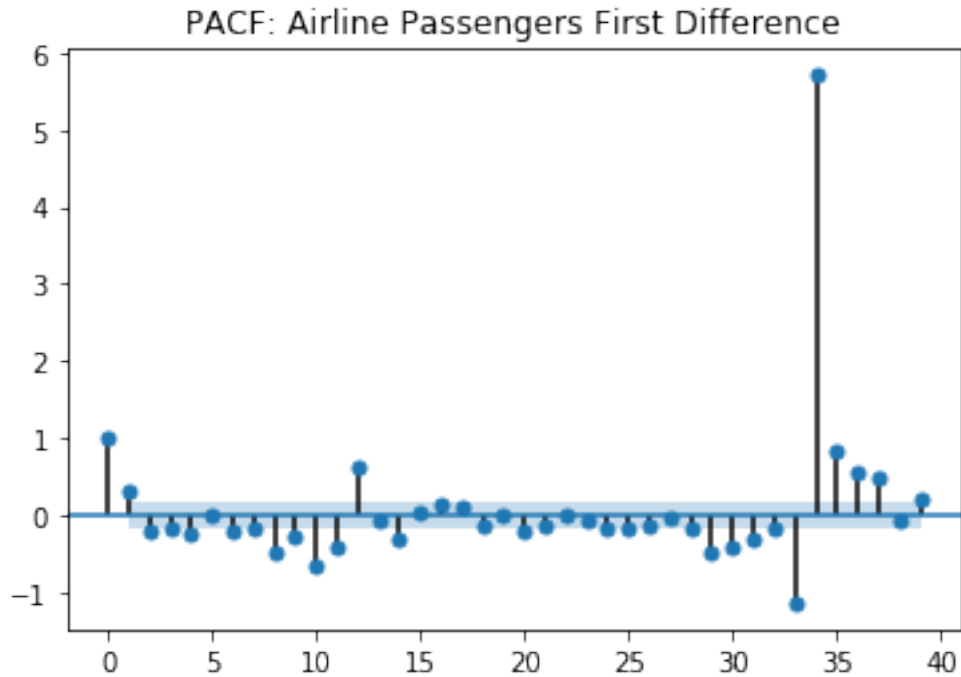
To make the Airline Passengers data stationary, we'll first apply differencing:

```
[18]: from statsmodels.tsa.statespace.tools import diff

df1['d1'] = diff(df1['Thousands of Passengers'],k_diff=1)
df1['d1'].plot(figsize=(12,5));
```



```
[19]: title='PACF: Airline Passengers First Difference'
lags=40
plot_pacf(df1['d1'].dropna(),title=title,lags=np.arange(lags)); # be sure to
↪add .dropna() here!
```




3.2.1 Plot Resizing

In case you want to display the full autocorrelation plot, it helps to increase the figure size using matplotlib:

```
[21]: import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize=(12,5))

plot_acf(df1['Thousands of Passengers'],ax=ax);
```



A NOTE ABOUT AUTOCORRELATION: Some texts compute lagged correlations using the Pearson Correlation Coefficient given by: $r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$

These are easily calculated in numpy with `numpy.corrcoef(x,y)` and in Excel with `=CORREL(x,y)`. Using our example, r_0 is still 1, but to solve for r_1 :

$$x_1 = [13, 5, 11, 12], \bar{x}_1 = 10.25 \quad y_1 = [5, 11, 12, 9], \bar{y}_1 = 9.25 \quad r_{x_1 y_1} = \frac{(13-10.25)(5-9.25) + (5-10.25)(11-9.25) + (11-10.25)(12-9.25) + (12-10.25)(9-9.25)}{\sqrt{((13-10.25)^2 + (5-10.25)^2 + (11-10.25)^2 + (12-10.25)^2)} \sqrt{((5-9.25)^2 + (11-9.25)^2 + (12-9.25)^2 + (9-9.25)^2)}} = \frac{-19.25}{33.38} = -0.577$$

However, there are some shortcomings. Using the Pearson method, the second-to-last term r_{k-1} will always be 1 and the last term r_k will always be undefined.