

04-Choosing-ARIMA-Orders

October 19, 2022

Copyright Pierian Data

For more information, visit us at www.pieriandata.com

1 Choosing ARIMA Orders

- Goals
- Understand PDQ terms for ARIMA (slides)
- Understand how to choose orders manually from ACF and PACF
- Understand how to use automatic order selection techniques using the functions below

Before we can apply an ARIMA forecasting model, we need to review the components of one. ARIMA, or Autoregressive Independent Moving Average is actually a combination of 3 models: * AR(p) Autoregression - a regression model that utilizes the dependent relationship between a current observation and observations over a previous period. * I(d) Integration - uses differencing of observations (subtracting an observation from an observation at the previous time step) in order to make the time series stationary * MA(q) Moving Average - a model that uses the dependency between an observation and a residual error from a moving average model applied to lagged observations.

Related Functions:

`pmdarima.auto_arima(y[,start_p,d,start_q, ...])` Returns the optimal order for an ARIMA model

Optional Function (see note below):

`stattools.arma_order_select_ic(y[, max_ar, ...])` Returns information criteria for many ARMA models
`x13.x13_arima_select_order(endog[, ...])` Perform automatic seasonal ARIMA order identification using x12/x13 ARIMA

1.1 Perform standard imports and load datasets

```
[1]: import pandas as pd
import numpy as np
%matplotlib inline
```

```

# Load a non-stationary dataset
df1 = pd.read_csv('../Data/airline_passengers.
    ↳csv', index_col='Month', parse_dates=True)
df1.index.freq = 'MS'

# Load a stationary dataset
df2 = pd.read_csv('../Data/DailyTotalFemaleBirths.
    ↳csv', index_col='Date', parse_dates=True)
df2.index.freq = 'D'

```

1.2 pmdarima Auto-ARIMA

This is a third-party tool separate from statsmodels. It should already be installed if you're using our virtual environment. If not, then at a terminal run: `pip install pmdarima`

```

[2]: from pmdarima import auto_arima

# Ignore harmless warnings
import warnings
warnings.filterwarnings("ignore")

```

```

[3]: help(auto_arima)

```

Help on function auto_arima in module pmdarima.arima.auto:

```

auto_arima(y, exogenous=None, start_p=2, d=None, start_q=2, max_p=5, max_d=2,
max_q=5, start_P=1, D=None, start_Q=1, max_P=2, max_D=1, max_Q=2, max_order=10,
m=1, seasonal=True, stationary=False, information_criterion='aic', alpha=0.05,
test='kpss', seasonal_test='ch', stepwise=True, n_jobs=1, start_params=None,
trend=None, method=None, transparams=True, solver='lbfgs', maxiter=50, disp=0,
callback=None, offset_test_args=None, seasonal_test_args=None,
suppress_warnings=False, error_action='warn', trace=False, random=False,
random_state=None, n_fits=10, return_valid_fits=False, out_of_sample_size=0,
scoring='mse', scoring_args=None, with_intercept=True, **fit_args)

```

Automatically discover the optimal order for an ARIMA model.

The ``auto_arima`` function seeks to identify the most optimal parameters for an ``ARIMA`` model, and returns a fitted ARIMA model. This function is based on the commonly-used R function, ``forecast::auto.arima`` [3].

The ``auto_arima`` function works by conducting differencing tests (i.e., Kwiatkowski-Phillips-Schmidt-Shin, Augmented Dickey-Fuller or Phillips-Perron) to determine the order of differencing, ``d``, and then fitting models within ranges of defined ``start_p``, ``max_p``, ``start_q``, ``max_q`` ranges. If the ``seasonal`` optional is enabled,

`auto_arima` also seeks to identify the optimal `P` and `Q` hyper-parameters after conducting the Canova-Hansen to determine the optimal order of seasonal differencing, `D`.

In order to find the best model, `auto_arima` optimizes for a given `information_criterion`, one of {'aic', 'aicc', 'bic', 'hqic', 'oob'} (Akaike Information Criterion, Corrected Akaike Information Criterion, Bayesian Information Criterion, Hannan-Quinn Information Criterion, or "out of bag"--for validation scoring--respectively) and returns the ARIMA which minimizes the value.

Note that due to stationarity issues, `auto_arima` might not find a suitable model that will converge. If this is the case, a `ValueError` will be thrown suggesting stationarity-inducing measures be taken prior to re-fitting or that a new range of `order` values be selected. Non-stepwise (i.e., essentially a grid search) selection can be slow, especially for seasonal data. Stepwise algorithm is outlined in Hyndman and Khandakar (2008).

Parameters

`y` : array-like or iterable, shape=(n_samples,)

The time-series to which to fit the `ARIMA` estimator. This may either be a Pandas `Series` object (statsmodels can internally use the dates in the index), or a numpy array. This should be a one-dimensional array of floats, and should not contain any `np.nan` or `np.inf` values.

`exogenous` : array-like, shape=[n_obs, n_vars], optional (default=None)

An optional 2-d array of exogenous variables. If provided, these variables are used as additional features in the regression operation. This should not include a constant or trend. Note that if an `ARIMA` is fit on exogenous features, it must be provided exogenous features for making predictions.

`start_p` : int, optional (default=2)

The starting value of `p`, the order (or number of time lags) of the auto-regressive ("AR") model. Must be a positive integer.

`d` : int, optional (default=None)

The order of first-differencing. If None (by default), the value will automatically be selected based on the results of the `test` (i.e., either the Kwiatkowski-Phillips-Schmidt-Shin, Augmented Dickey-Fuller or the Phillips-Perron test will be conducted to find the most probable value). Must be a positive integer or None. Note that if `d` is None, the runtime could be significantly longer.

`start_q` : int, optional (default=2)

The starting value of ``q``, the order of the moving-average ("MA") model. Must be a positive integer.

max_p : int, optional (default=5)

The maximum value of ``p``, inclusive. Must be a positive integer greater than or equal to ``start_p``.

max_d : int, optional (default=2)

The maximum value of ``d``, or the maximum number of non-seasonal differences. Must be a positive integer greater than or equal to ``d``.

max_q : int, optional (default=5)

The maximum value of ``q``, inclusive. Must be a positive integer greater than ``start_q``.

start_P : int, optional (default=1)

The starting value of ``P``, the order of the auto-regressive portion of the seasonal model.

D : int, optional (default=None)

The order of the seasonal differencing. If None (by default, the value will automatically be selected based on the results of the ``seasonal_test``. Must be a positive integer or None.

start_Q : int, optional (default=1)

The starting value of ``Q``, the order of the moving-average portion of the seasonal model.

max_P : int, optional (default=2)

The maximum value of ``P``, inclusive. Must be a positive integer greater than ``start_P``.

max_D : int, optional (default=1)

The maximum value of ``D``. Must be a positive integer greater than ``D``.

max_Q : int, optional (default=2)

The maximum value of ``Q``, inclusive. Must be a positive integer greater than ``start_Q``.

max_order : int, optional (default=10)

If the sum of ``p`` and ``q`` is \geq ``max_order``, a model will *not* be fit with those parameters, but will progress to the next combination. Default is 5. If ``max_order`` is None, it means there are no constraints on maximum order.

m : int, optional (default=1)

The period for seasonal differencing, ``m`` refers to the number of

periods in each season. For example, ``m`` is 4 for quarterly data, 12 for monthly data, or 1 for annual (non-seasonal) data. Default is 1. Note that if ``m`` == 1 (i.e., is non-seasonal), ``seasonal`` will be set to False. For more information on setting this parameter, see :ref:`period`.

`seasonal` : bool, optional (default=True)
Whether to fit a seasonal ARIMA. Default is True. Note that if ``seasonal`` is True and ``m`` == 1, ``seasonal`` will be set to False.

`stationary` : bool, optional (default=False)
Whether the time-series is stationary and ``d`` should be set to zero.

`information_criterion` : str, optional (default='aic')
The information criterion used to select the best ARIMA model. One of ``pmdarima.arima.auto_arima.VALID_CRITERIA``, ('aic', 'bic', 'hqic', 'oob').

`alpha` : float, optional (default=0.05)
Level of the test for testing significance.

`test` : str, optional (default='kpss')
Type of unit root test to use in order to detect stationarity if ``stationary`` is False and ``d`` is None. Default is 'kpss' (Kwiatkowski-Phillips-Schmidt-Shin).

`seasonal_test` : str, optional (default='ch')
This determines which seasonal unit root test is used if ``seasonal`` is True and ``D`` is None. Default is 'ch' (Canova-Hansen).

`stepwise` : bool, optional (default=True)
Whether to use the stepwise algorithm outlined in Hyndman and Khandakar (2008) to identify the optimal model parameters. The stepwise algorithm can be significantly faster than fitting all (or a ``random`` subset of) hyper-parameter combinations and is less likely to over-fit the model.

`n_jobs` : int, optional (default=1)
The number of models to fit in parallel in the case of a grid search (``stepwise=False``). Default is 1, but -1 can be used to designate "as many as possible".

`start_params` : array-like, optional (default=None)
Starting parameters for ``ARMA(p,q)``. If None, the default is given by ``ARMA._fit_start_params``.

`transparams` : bool, optional (default=True)
Whether or not to transform the parameters to ensure stationarity.

Uses the transformation suggested in Jones (1980). If False, no checking for stationarity or invertibility is done.

method : str, one of {'css-mle', 'mle', 'css'}, optional (default=None)

This is the loglikelihood to maximize. If "css-mle", the conditional sum of squares likelihood is maximized and its values are used as starting values for the computation of the exact likelihood via the Kalman filter. If "mle", the exact likelihood is maximized via the Kalman Filter. If "css" the conditional sum of squares likelihood is maximized. All three methods use `start_params` as starting parameters. See above for more information. If fitting a seasonal ARIMA, the default is 'lbfgs'

trend : str or None, optional (default=None)

The trend parameter. If ``with_intercept`` is True, ``trend`` will be used. If ``with_intercept`` is False, the trend will be set to a no-intercept value.

solver : str or None, optional (default='lbfgs')

Solver to be used. The default is 'lbfgs' (limited memory Broyden-Fletcher-Goldfarb-Shanno). Other choices are 'bfgs', 'newton' (Newton-Raphson), 'nm' (Nelder-Mead), 'cg' - (conjugate gradient), 'ncg' (non-conjugate gradient), and 'powell'. By default, the limited memory BFGS uses m=12 to approximate the Hessian, projected gradient tolerance of 1e-8 and factr = 1e2. You can change these by using kwargs.

maxiter : int, optional (default=50)

The maximum number of function evaluations. Default is 50.

disp : int, optional (default=0)

If True, convergence information is printed. For the default 'lbfgs' ``solver``, disp controls the frequency of the output during the iterations. disp < 0 means no output in this case.

callback : callable, optional (default=None)

Called after each iteration as callback(xk) where xk is the current parameter vector. This is only used in non-seasonal ARIMA models.

offset_test_args : dict, optional (default=None)

The args to pass to the constructor of the offset (``d``) test. See ``pmdarima.arima.stationarity`` for more details.

seasonal_test_args : dict, optional (default=None)

The args to pass to the constructor of the seasonal offset (``D``) test. See ``pmdarima.arima.seasonality`` for more details.

suppress_warnings : bool, optional (default=False)

Many warnings might be thrown inside of statsmodels. If `suppress_warnings` is True, all of the warnings coming from ARIMA` will be squelched.`

`error_action` : str, optional (default='warn')

If unable to fit an `ARIMA` due to stationarity issues, whether to warn ('warn'), raise the ValueError` ('raise') or ignore ('ignore'). Note that the default behavior is to warn, and fits that fail will be returned as None. This is the recommended behavior, as statsmodels ARIMA and SARIMAX models hit bugs periodically that can cause an otherwise healthy parameter combination to fail for reasons not related to pmdarima.`

`trace` : bool, optional (default=False)

Whether to print status on the fits. Note that this can be very verbose...

`random` : bool, optional (default=False)

Similar to grid searches, `auto_arima` provides the capability to perform a "random search" over a hyper-parameter space. If random` is True, rather than perform an exhaustive search or stepwise` search, only n_fits` ARIMA models will be fit (stepwise` must be False for this option to do anything).`

`random_state` : int, long or numpy `RandomState`, optional (default=None)`

The PRNG for when `random=True`. Ensures replicable testing and results.`

`n_fits` : int, optional (default=10)

If `random` is True and a "random search" is going to be performed, n_iter` is the number of ARIMA models to be fit.`

`return_valid_fits` : bool, optional (default=False)

If True, will return all valid ARIMA fits in a list. If False (by default), will only return the best fit.

`out_of_sample_size` : int, optional (default=0)

The `ARIMA` class can fit only a portion of the data if specified, in order to retain an "out of bag" sample score. This is the number of examples from the tail of the time series to hold out and use as validation examples. The model will not be fit on these samples, but the observations will be added into the model's endog` and exog` arrays so that future forecast values originate from the end of the endogenous vector.`

For instance::

```
y = [0, 1, 2, 3, 4, 5, 6]
```

```

    out_of_sample_size = 2

    > Fit on: [0, 1, 2, 3, 4]
    > Score on: [5, 6]
    > Append [5, 6] to end of self.arima_res_.data.endog values

scoring : str, optional (default='mse')
    If performing validation (i.e., if ``out_of_sample_size`` > 0), the
    metric to use for scoring the out-of-sample data. One of {'mse', 'mae'}

scoring_args : dict, optional (default=None)
    A dictionary of key-word arguments to be passed to the ``scoring``
    metric.

with_intercept : bool, optional (default=True)
    Whether to include an intercept term. Default is True.

**fit_args : dict, optional (default=None)
    A dictionary of keyword arguments to pass to the :func:`ARIMA.fit`
    method.

See Also
-----
:func:`pmdarima.arima.ARIMA`

Notes
-----
Fitting with `stepwise=False` can prove slower, especially when
`seasonal=True`.

References
-----
.. [1] https://wikipedia.org/wiki/Autoregressive\_integrated\_moving\_average
.. [2] R's auto-arima source code: http://bit.ly/2g0h5z2
.. [3] R's auto-arima documentation: http://bit.ly/2wbBvUN

```

Let's look first at the stationary, non-seasonal Daily Female Births dataset:

```
[4]: auto_arima(df2['Births'])
```

```
[4]: ARIMA(callback=None, disp=0, maxiter=50, method=None, order=(1, 1, 1),
    out_of_sample_size=0, scoring='mse', scoring_args={},
    seasonal_order=(0, 0, 0, 1), solver='lbfgs', start_params=None,
    suppress_warnings=False, transparams=True, trend=None,
    with_intercept=True)
```

NOTE: Harmless warnings should have been suppressed, but if you see an error citing unusual

behavior you can suppress this message by passing `error_action='ignore'` into `auto_arma()`. Also, `auto_arma().summary()` provides a nicely formatted summary table.

```
[5]: auto_arma(df2['Births'], error_action='ignore').summary()
```

```
[5]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                Statespace Model Results
      =====
Dep. Variable:                  y      No. Observations:                  365
Model:                        SARIMAX(1, 1, 1)  Log Likelihood                -1226.077
Date:                        Fri, 22 Mar 2019  AIC                      2460.154
Time:                        10:46:05      BIC                      2475.743
Sample:                      0      HQIC                      2466.350
                                - 365
Covariance Type:              opg
      =====
                                coef      std err          z      P>|z|      [0.025      0.975]
      -----
intercept          0.0132      0.014      0.975      0.330      -0.013      0.040
ar.L1              0.1299      0.059      2.217      0.027      0.015      0.245
ma.L1             -0.9694      0.016     -62.235      0.000     -1.000     -0.939
sigma2            48.9989      3.432     14.279      0.000     42.273     55.725
      =====
      ==
Ljung-Box (Q):                36.69      Jarque-Bera (JB):
26.17
Prob(Q):                      0.62      Prob(JB):
0.00
Heteroskedasticity (H):        0.97      Skew:
0.58
Prob(H) (two-sided):          0.85      Kurtosis:
3.62
      =====
      ==

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
      """
```

This shows a recommended (p,d,q) ARIMA Order of (1,1,1), with no seasonal_order component.

We can see how this was determined by looking at the stepwise results. The recommended order is the one with the lowest Akaike information criterion or AIC score. Note that the recommended model may not be the one with the closest fit. The AIC score takes complexity into account, and tries to identify the best forecasting model.

```
[6]: stepwise_fit = auto_arima(df2['Births'], start_p=0, start_q=0,
                             max_p=6, max_q=3, m=12,
                             seasonal=False,
                             d=None, trace=True,
                             error_action='ignore', # we don't want to know if
→an order does not work
                             suppress_warnings=True, # we don't want convergence
→warnings
                             stepwise=True) # set to stepwise

stepwise_fit.summary()
```

```
Fit ARIMA: order=(0, 1, 0); AIC=2650.760, BIC=2658.555, Fit time=0.018 seconds
Fit ARIMA: order=(1, 1, 0); AIC=2565.234, BIC=2576.925, Fit time=0.139 seconds
Fit ARIMA: order=(0, 1, 1); AIC=2463.584, BIC=2475.275, Fit time=0.047 seconds
Fit ARIMA: order=(1, 1, 1); AIC=2460.154, BIC=2475.742, Fit time=0.105 seconds
Fit ARIMA: order=(1, 1, 2); AIC=2460.515, BIC=2480.000, Fit time=0.320 seconds
Fit ARIMA: order=(2, 1, 2); AIC=2462.045, BIC=2485.428, Fit time=0.373 seconds
Fit ARIMA: order=(2, 1, 1); AIC=2461.271, BIC=2480.757, Fit time=0.169 seconds
Total fit time: 1.175 seconds
```

```
[6]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                ARIMA Model Results
=====
Dep. Variable:                  D.y    No. Observations:                  364
Model:                        ARIMA(1, 1, 1)    Log Likelihood                  -1226.077
Method:                        css-mle    S.D. of innovations                  7.000
Date:                          Fri, 22 Mar 2019    AIC                  2460.154
Time:                          10:46:20    BIC                  2475.742
Sample:                        1    HQIC                  2466.350
=====

              coef      std err          z      P>|z|      [0.025      0.975]
-----
const          0.0152      0.014        1.068      0.286      -0.013      0.043
ar.L1.D.y       0.1299      0.056        2.334      0.020       0.021      0.239
ma.L1.D.y      -0.9694      0.019     -51.415      0.000      -1.006     -0.932
              Roots
=====
              Real      Imaginary      Modulus      Frequency
-----
AR.1          7.6996      +0.0000j        7.6996      0.0000
MA.1          1.0316      +0.0000j        1.0316      0.0000
=====
"""
```

Now let's look at the non-stationary, seasonal Airline Passengers dataset:

```
[7]: stepwise_fit = auto_arma(df1['Thousands of Passengers'], start_p=1, start_q=1,
                             max_p=3, max_q=3, m=12,
                             start_P=0, seasonal=True,
                             d=None, D=1, trace=True,
                             error_action='ignore', # we don't want to know if
                             ↪an order does not work
                             suppress_warnings=True, # we don't want convergence
                             ↪warnings
                             stepwise=True)          # set to stepwise

stepwise_fit.summary()
```

```
Fit ARIMA: order=(1, 1, 1) seasonal_order=(0, 1, 1, 12); AIC=1024.824,
BIC=1039.200, Fit time=0.574 seconds
Fit ARIMA: order=(0, 1, 0) seasonal_order=(0, 1, 0, 12); AIC=1033.479,
BIC=1039.229, Fit time=0.000 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(1, 1, 0, 12); AIC=1022.316,
BIC=1033.817, Fit time=0.294 seconds
Fit ARIMA: order=(0, 1, 1) seasonal_order=(0, 1, 1, 12); AIC=1022.904,
BIC=1034.405, Fit time=0.251 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(0, 1, 0, 12); AIC=1022.343,
BIC=1030.968, Fit time=0.094 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(2, 1, 0, 12); AIC=1021.142,
BIC=1035.518, Fit time=0.748 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(2, 1, 1, 12); AIC=1016.960,
BIC=1034.211, Fit time=2.258 seconds
Fit ARIMA: order=(0, 1, 0) seasonal_order=(2, 1, 1, 12); AIC=1033.371,
BIC=1047.747, Fit time=1.931 seconds
Fit ARIMA: order=(2, 1, 0) seasonal_order=(2, 1, 1, 12); AIC=1018.094,
BIC=1038.221, Fit time=2.492 seconds
Fit ARIMA: order=(1, 1, 1) seasonal_order=(2, 1, 1, 12); AIC=1017.829,
BIC=1037.955, Fit time=2.464 seconds
Fit ARIMA: order=(2, 1, 1) seasonal_order=(2, 1, 1, 12); AIC=nan, BIC=nan, Fit
time=nan seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(1, 1, 1, 12); AIC=1022.425,
BIC=1036.801, Fit time=0.432 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(2, 1, 2, 12); AIC=1017.410,
BIC=1037.536, Fit time=2.546 seconds
Total fit time: 14.093 seconds
```

```
[7]: <class 'statsmodels.iolib.summary.Summary'>
     """
```

Statespace Model Results

=====

```

=====
Dep. Variable:          y    No. Observations:
144
Model:                SARIMAX(1, 1, 0)x(2, 1, 1, 12)    Log Likelihood
-502.480
Date:                  Fri, 22 Mar 2019    AIC
1016.960
Time:                  10:46:39    BIC
1034.211
Sample:                0    HQIC
1023.970

                                - 144
Covariance Type:        opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
intercept	0.0045	0.178	0.025	0.980	-0.345	0.354
ar.L1	-0.3766	0.077	-4.890	0.000	-0.527	-0.226
ar.S.L12	0.6891	0.140	4.918	0.000	0.414	0.964
ar.S.L24	0.3091	0.107	2.883	0.004	0.099	0.519
ma.S.L12	-0.9742	0.511	-1.906	0.057	-1.976	0.028
sigma2	113.2075	48.855	2.317	0.020	17.453	208.961

```

=====
===
Ljung-Box (Q):          58.67    Jarque-Bera (JB):
12.12
Prob(Q):                0.03    Prob(JB):
0.00
Heteroskedasticity (H): 2.70    Skew:
0.10
Prob(H) (two-sided):    0.00    Kurtosis:
4.48
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""

```

1.3 OPTIONAL: statsmodels ARMA_Order_Select_IC

Statsmodels has a selection tool to find orders for ARMA models on stationary data.

```
[8]: from statsmodels.tsa.stattools import arma_order_select_ic
```

```
[9]: help(arma_order_select_ic)
```

Help on function arma_order_select_ic in module statsmodels.tsa.stattools:

```
arma_order_select_ic(y, max_ar=4, max_ma=2, ic='bic', trend='c', model_kw={},
fit_kw={})
```

Returns information criteria for many ARMA models

Parameters

y : array-like

Time-series data

max_ar : int

Maximum number of AR lags to use. Default 4.

max_ma : int

Maximum number of MA lags to use. Default 2.

ic : str, list

Information criteria to report. Either a single string or a list of different criteria is possible.

trend : str

The trend to use when fitting the ARMA models.

model_kw : dict

Keyword arguments to be passed to the ``ARMA`` model

fit_kw : dict

Keyword arguments to be passed to ``ARMA.fit``.

Returns

obj : Results object

Each ic is an attribute with a DataFrame for the results. The AR order used is the row index. The ma order used is the column index. The minimum orders are available as ``ic_min_order``.

Examples

```
>>> from statsmodels.tsa.arima_process import arma_generate_sample
```

```
>>> import statsmodels.api as sm
```

```
>>> import numpy as np
```

```
>>> arparams = np.array([.75, -.25])
```

```
>>> maparams = np.array([.65, .35])
```

```
>>> arparams = np.r_[1, -arparams]
```

```
>>> maparam = np.r_[1, maparams]
```

```
>>> nobs = 250
```

```
>>> np.random.seed(2014)
```

```
>>> y = arma_generate_sample(arparams, maparams, nobs)
```

```
>>> res = sm.tsa.arma_order_select_ic(y, ic=['aic', 'bic'], trend='nc')
>>> res.aic_min_order
>>> res.bic_min_order
```

Notes

This method can be used to tentatively identify the order of an ARMA process, provided that the time series is stationary and invertible. This function computes the full exact MLE estimate of each model and can be, therefore a little slow. An implementation using approximate estimates will be provided in the future. In the meantime, consider passing {method : 'css'} to fit_kw.

```
[10]: arma_order_select_ic(df2['Births'])
```

```
[10]: {'bic':
      0      1      2
0  2502.581666  2494.238827  2494.731525
1  2490.780306  2484.505386      NaN
2  2491.963234  2485.782753  2491.097206
3  2496.498618  2491.061564  2496.961178
4  2501.491891  2504.012579  2498.329743, 'bic_min_order': (1, 1)}
```

```
[11]: arma_order_select_ic(df1['Thousands of Passengers'])
```

```
[11]: {'bic':
      0      1      2
0  1796.307207  1627.771967  1534.002384
1  1437.088819  1421.627524  1425.899321
2  1425.518037  1423.098290      NaN
3  1425.191373      NaN  1400.744437
4  1427.576572      NaN  1414.310652, 'bic_min_order': (3, 2)}
```

A note about statsmodels.tsa.x13.x13_arima_select_order This utility requires installation of X-13ARIMA-SEATS, a seasonal adjustment tool developed by the U.S. Census Bureau. See <https://www.census.gov/srd/www/x13as/> for details. Since the installation requires adding x13as to your PATH settings we've deemed it beyond the scope of this course.