

10-Forecasting-Exercises

October 19, 2022

Copyright Pierian Data

For more information, visit us at www.pieriandata.com

1 Forecasting Exercises

This exercise walks through a SARIMA prediction and forecast similar to the one done on the Mauna Loa CO₂ dataset. This time we're using a seasonal time series of California Hospitality Industry Employees.

IMPORTANT NOTE! Make sure you don't run the cells directly above the example output shown, otherwise you will end up writing over the example output!

```
[1]: # RUN THIS CELL
import pandas as pd
import numpy as np
%matplotlib inline

# Load specific forecasting tools
from statsmodels.tsa.statespace.sarimax import SARIMAX

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf # for determining
    ↪ (p, q) orders
from statsmodels.tsa.seasonal import seasonal_decompose      # for ETS Plots
from pmdarima import auto_arima                             # for determining
    ↪ ARIMA orders

# Load specific evaluation tools
from sklearn.metrics import mean_squared_error
from statsmodels.tools.eval_measures import rmse

# Ignore harmless warnings
import warnings
warnings.filterwarnings("ignore")
```

```
# Load datasets
df = pd.read_csv('../Data/HospitalityEmployees.
→csv',index_col='Date',parse_dates=True)
df.index.freq = 'MS'
print(len(df))
print(df.head())
```

348

Date	Employees
1990-01-01	1064.5
1990-02-01	1074.5
1990-03-01	1090.0
1990-04-01	1097.4
1990-05-01	1108.7


So df has 348 records and one column. The data represents the number of employees in thousands of persons as monthly averages from January, 1990 to December 2018.

1.0.1 1. Plot the source data

Create a line chart of the dataset. Optional: add a title and y-axis label.

```
[ ]: ## CODE HERE
```

```
[2]: # DON'T WRITE HERE
```




1.0.2 2. Run an ETS Decomposition

Use an 'additive' model.

[]:

[3]: *# DON'T WRITE HERE*



1.0.3 3. Run `pmdarima.auto_arima` to obtain recommended orders

This may take awhile as there are a lot of combinations to evaluate.

[]:

[4]: *# DON'T WRITE HERE*

[4]: `<class 'statsmodels.iolib.summary.Summary'>`

"""

Statespace Model Results

=====

Dep. Variable: y No. Observations:
348

Model: SARIMAX(0, 1, 0)x(2, 0, 0, 12) Log Likelihood
-1134.664

Date: Wed, 27 Mar 2019 AIC
2277.328

```
Time: 13:12:10 BIC
2292.726
Sample: 0 HQIC
2283.459
```

- 348

```
Covariance Type: opg
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-0.0477	0.292	-0.163	0.870	-0.620	0.524
ar.S.L12	0.5291	0.040	13.286	0.000	0.451	0.607
ar.S.L24	0.4303	0.041	10.453	0.000	0.350	0.511
sigma2	37.2952	2.157	17.294	0.000	33.068	41.522

===

```
Ljung-Box (Q): 99.53 Jarque-Bera (JB):
51.67
Prob(Q): 0.00 Prob(JB):
0.00
Heteroskedasticity (H): 0.86 Skew:
-0.29
Prob(H) (two-sided): 0.42 Kurtosis:
4.80
```

=====

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""
```

You should see a recommended ARIMA Order of (0,1,0) combined with a seasonal order of (2,0,0,12). ### 4. Split the data into train/test sets Set one year (12 records) for testing. There is more than one way to do this!

```
[ ]:
```

```
[5]: # DON'T WRITE HERE
```

1.0.4 5. Fit a SARIMA(0,1,0)(2,0,0,12) model to the training set

```
[ ]:
```

```
[6]: # DON'T WRITE HERE
```

```
[6]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                Statespace Model Results
=====
=====
Dep. Variable:                    Employees    No. Observations:
336
Model:                          SARIMAX(0, 1, 0)x(2, 0, 0, 12)    Log Likelihood
-1095.407
Date:                            Wed, 27 Mar 2019    AIC
2196.814
Time:                            13:12:24    BIC
2208.256
Sample:                          01-01-1990    HQIC
2201.375
                                - 12-01-2017
Covariance Type:                  opg
=====
                                coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.S.L12          0.5204      0.040      13.051      0.000      0.442      0.599
ar.S.L24          0.4385      0.041      10.593      0.000      0.357      0.520
sigma2           37.1907      2.165      17.175      0.000     32.947     41.435
=====
===
Ljung-Box (Q):                102.80    Jarque-Bera (JB):
56.66
Prob(Q):                      0.00    Prob(JB):
0.00
Heteroskedasticity (H):       1.06    Skew:
-0.35
Prob(H) (two-sided):          0.74    Kurtosis:
4.89
=====
===

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""
```

1.0.5 6. Obtain predicted values

```
[ ]:
```


```
[7]: # DON'T WRITE HERE
```

1.0.6 7. Plot predictions against known values

Optional: add a title and y-axis label.

```
[ ]:
```

```
[8]: # DON'T WRITE HERE
```



1.0.7 8. Evaluate the Model using MSE and RMSE

You can run both from the same cell if you want.

```
[ ]:
```

```
[9]: # DON'T WRITE HERE
```

```
SARIMA(0,1,0)(2,0,0,12) MSE Error: 182.8506646  
SARIMA(0,1,0)(2,0,0,12) RMSE Error: 13.52222854
```

1.0.8 9. Retrain the model on the full data and forecast one year into the future

```
[ ]:
```

```
[10]: # DON'T WRITE HERE
```

1.0.9 10. Plot the forecasted values alongside the original data

Optional: add a title and y-axis label.

```
[ ]:
```

```
[11]: # DON'T WRITE HERE
```



1.1 Great job!