

03-RNN-Exercise-Solutions

October 19, 2022

Copyright Pierian Data

For more information, visit us at www.pieriandata.com

1 RNN Example for Time Series

TASK: IMPORT THE BASIC LIBRARIES YOU THINK YOU WILL USE

```
[1]: import pandas as pd
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt
```

1.1 Data

Info about this data set: <https://fred.stlouisfed.org/series/TRFVOLUSM227NFWA>

Read in the data set "Miles_Traveled.csv" from the Data folder. Figure out how to set the date to a datetime index columns

```
[2]: # CODE HERE
```

```
[3]: df = pd.read_csv('../Data/Miles_Traveled.csv', index_col='DATE', parse_dates=True)
df.index.freq = 'MS'
```

```
[4]: df.head()
```

```
[4]:          TRFVOLUSM227NFWA
DATE
1970-01-01          80173.0
1970-02-01          77442.0
1970-03-01          90223.0
1970-04-01          89956.0
1970-05-01          97972.0
```


Task: Change the column names to Value

```
[5]: df.columns = ['Value']
```

TASK: Plot out the time series

```
[6]: df.plot(figsize=(12,8))
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x27987751240>
```




TASK: Perform a Seasonal Decomposition on the model and plot out the ETS components

```
[7]: # CODE HERE
```

```
[8]: from statsmodels.tsa.seasonal import seasonal_decompose
```


```
[9]: results = seasonal_decompose(df['Value'])  
     results.observed.plot(figsize=(12,2))
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x279898815c0>
```




```
[10]: results.trend.plot(figsize=(12,2))
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x27987a4eac8>
```




```
[11]: results.seasonal.plot(figsize=(12,2))
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x279899b64e0>
```



```
[12]: results.resid.plot(figsize=(12,2))
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x279899a9828>
```



1.2 Train Test Split

TASK: Figure out the length of the data set

```
[13]: len(df)
```

```
[13]: 588
```

```
[14]: train_len = len(df)-12
```

TASK: Split the data into a train/test split where the test set is the last 12 months of data.

```
[15]: train = df.iloc[:train_len]  
test = df.iloc[train_len:]
```

```
[16]: len(test)
```

```
[16]: 12
```

1.3 Scale Data

TASK: Use a MinMaxScaler to scale the train and test sets into scaled versions.

```
[17]: # CODE HERE
```

```
[18]: from sklearn.preprocessing import MinMaxScaler
```

```
[19]: scaler = MinMaxScaler()
```

```
[20]: # IGNORE WARNING ITS JUST CONVERTING TO FLOATS  
# WE ONLY FIT TO TRAININ DATA, OTHERWISE WE ARE CHEATING ASSUMING INFO ABOUT_  
#   ↳TEST SET  
scaler.fit(train)
```

```
[20]: MinMaxScaler(copy=True, feature_range=(0, 1))
```

```
[21]: scaled_train = scaler.transform(train)  
scaled_test = scaler.transform(test)
```

2 Time Series Generator

TASK: Create a TimeSeriesGenerator object based off the scaled_train data. The n_input is up to you, but at a minimum it should be at least 12.

```
[22]: #CODE HERE
```

```
[23]: from keras.preprocessing.sequence import TimeseriesGenerator
```

Using TensorFlow backend.

```
[24]: n_input = 24
      n_features=1
      generator = TimeseriesGenerator(scaled_train, scaled_train, length=n_input,
      ↪batch_size=1)
```

2.0.1 Create the Model

TASK: Create a Keras Sequential Model with as many LSTAM units you want and a final Dense Layer.

```
[25]: from keras.models import Sequential
      from keras.layers import Dense
      from keras.layers import LSTM
```

```
[26]: # define model
      model = Sequential()
      model.add(LSTM(150, activation='relu', input_shape=(n_input, n_features)))
      model.add(Dense(1))
      model.compile(optimizer='adam', loss='mse')
```

```
[27]: model.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
=====
lstm_1 (LSTM)                 (None, 150)              91200
-----
dense_1 (Dense)               (None, 1)                151
=====
Total params: 91,351
Trainable params: 91,351
Non-trainable params: 0
-----
```

TASK: Fit the model to the generator (it should be a lot of epochs, but do as many as you have the patience for! :)

```
[28]: # CODE HERE
```

```
[29]: # fit model  
model.fit_generator(generator, epochs=30)
```

```
Epoch 1/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0139  
Epoch 2/30  
552/552 [=====] - 17s 30ms/step - loss: 0.0070  
Epoch 3/30  
552/552 [=====] - 16s 30ms/step - loss: 0.0056  
Epoch 4/30  
552/552 [=====] - 16s 29ms/step - loss: 0.0050  
Epoch 5/30  
552/552 [=====] - 17s 30ms/step - loss: 0.0036  
Epoch 6/30  
552/552 [=====] - 16s 30ms/step - loss: 0.0025  
Epoch 7/30  
552/552 [=====] - 16s 30ms/step - loss: 0.0025  
Epoch 8/30  
552/552 [=====] - 17s 30ms/step - loss: 0.0018  
Epoch 9/30  
552/552 [=====] - 18s 32ms/step - loss: 0.0016  
Epoch 10/30  
552/552 [=====] - 19s 34ms/step - loss: 0.0017  
Epoch 11/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0016  
Epoch 12/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0013  
Epoch 13/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0014  
Epoch 14/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0012 0s - 1  
Epoch 15/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0013  
Epoch 16/30  
552/552 [=====] - 19s 34ms/step - loss: 0.0011  
Epoch 17/30  
552/552 [=====] - 17s 32ms/step - loss: 0.0011  
Epoch 18/30  
552/552 [=====] - 18s 33ms/step - loss: 9.6872e-04  
Epoch 19/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0011  
Epoch 20/30  
552/552 [=====] - 18s 33ms/step - loss: 0.0011  
Epoch 21/30  
552/552 [=====] - 18s 33ms/step - loss: 9.6653e-04
```

```

Epoch 22/30
552/552 [=====] - 18s 33ms/step - loss: 9.4804e-04
Epoch 23/30
552/552 [=====] - 18s 33ms/step - loss: 8.5220e-04
Epoch 24/30
552/552 [=====] - 18s 33ms/step - loss: 7.8977e-04
Epoch 25/30
552/552 [=====] - 18s 33ms/step - loss: 7.9855e-04
Epoch 26/30
552/552 [=====] - 18s 33ms/step - loss: 8.0307e-04
Epoch 27/30
552/552 [=====] - 18s 33ms/step - loss: 7.8170e-04
Epoch 28/30
552/552 [=====] - 18s 33ms/step - loss: 7.8733e-04
Epoch 29/30
552/552 [=====] - 18s 33ms/step - loss: 7.7737e-04
Epoch 30/30
552/552 [=====] - 18s 33ms/step - loss: 7.7677e-04

```

[29]: <keras.callbacks.History at 0x2798f3695f8>

TASK: Plot the history of the loss that occurred during training.


[30]: *# CODE HERE*

[31]: `model.history.history.keys()`

[31]: `dict_keys(['loss'])`

[32]: `loss_per_epoch = model.history.history['loss']
plt.plot(range(len(loss_per_epoch)), loss_per_epoch)`

[32]: [`<matplotlib.lines.Line2D at 0x27bf44532e8>`]



2.1 Evaluate on Test Data

TASK: Based on your test data and input size, create an appropriate;y sized "first evaluation batch" like we did in the lecture.

```
[33]: first_eval_batch = scaled_train[-12:]
```

```
[34]: first_eval_batch
```

```
[34]: array([[0.79630397],
             [0.71226435],
             [0.90477416],
             [0.93121043],
             [0.98386382],
             [0.96757519],
             [1.         ],
             [0.9801859 ],
             [0.8824684 ],
             [0.95995255],
             [0.85883345],
             [0.90086755]])
```

```
[35]: first_eval_batch = first_eval_batch.reshape((1, 12, n_features))
```

TASK: Generate predictions into the same time stamps as the test set


```
[36]: # CODE HERE
```

```
[37]: test_predictions = []

first_eval_batch = scaled_train[-n_input:]
current_batch = first_eval_batch.reshape((1, n_input, n_features))

for i in range(len(test)):

    # get prediction 1 time stamp ahead ([0] is for grabbing just the number,
    ↪ instead of [array])
    current_pred = model.predict(current_batch)[0]

    # store prediction
    test_predictions.append(current_pred)

    # update batch to now include prediction and drop first value
    current_batch = np.append(current_batch[:,1:,:], [[current_pred]], axis=1)
```

2.2 Inverse Transformations and Compare

TASK: Inverse Transform your new forecasted predictions.

```
[38]: #CODE HERE
```

```
[39]: true_predictions = scaler.inverse_transform(test_predictions)
```

```
[40]: true_predictions
```

```
[40]: array([[244790.4583751 ],
            [233844.90710074],
            [270457.81608635],
            [272201.29687929],
            [279289.6588499 ],
            [279346.60926121],
            [286210.98546427],
            [281885.08126384],
            [268614.50942963],
            [275487.78100669],
            [258639.72996229],
            [262481.15531212]])
```

TASK: Create a new dataframe that has both the original test values and your predictions for them.

```
[41]: # IGNORE WARNINGS
test['Predictions'] = true_predictions
```

```
C:\Users\Marcial\Anaconda3\lib\site-packages\ipykernel_launcher.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
[42]: test
```


```
[42]:
```

	Value	Predictions
DATE		
2018-01-01	245695.0	244790.458375
2018-02-01	226660.0	233844.907101
2018-03-01	268480.0	270457.816086
2018-04-01	272475.0	272201.296879
2018-05-01	286164.0	279289.658850
2018-06-01	280877.0	279346.609261
2018-07-01	288145.0	286210.985464
2018-08-01	286608.0	281885.081264
2018-09-01	260595.0	268614.509430
2018-10-01	282174.0	275487.781007
2018-11-01	258590.0	258639.729962
2018-12-01	268413.0	262481.155312

TASK: Plot out the test set against your own predicted values.

```
[43]: test.plot(figsize=(12,8))
```

```
[43]: <matplotlib.axes._subplots.AxesSubplot at 0x279bcda4588>
```



3 Saving Models

TASK: Optional, Save your model!

```
[44]: model.save('solutions_model.h5')
```